

Self-Driving Car Engineer Nanodegree

Deep Learning

Project: Build a Traffic Sign Recognition Classifier

In this notebook, a template is provided for you to implement your functionality in stages, which is required to successfully complete this project. If additional code is required that cannot be included in the notebook, be sure that the Python code is successfully imported and included in your submission if necessary.

Note: Once you have completed all of the code implementations, you need to finalize your work by exporting the iPython Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to "\n", "**File -> Download as -> HTML (.html)**". Include the finished document along with this notebook as your submission.

In addition to implementing code, there is a writeup to complete. The writeup should be completed in a separate file, which can be either a markdown file or a pdf document. There is a [write up template \(https://github.com/udacity/CarND-Traffic-Sign-Classifier-Project/blob/master/writeup_template.md\)](https://github.com/udacity/CarND-Traffic-Sign-Classifier-Project/blob/master/writeup_template.md) that can be used to guide the writing process. Completing the code template and writeup template will cover all of the [rubric points \(https://review.udacity.com/#!/rubrics/481/view\)](https://review.udacity.com/#!/rubrics/481/view) for this project.

The [rubric \(https://review.udacity.com/#!/rubrics/481/view\)](https://review.udacity.com/#!/rubrics/481/view) contains "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. The stand out suggestions are optional. If you decide to pursue the "stand out suggestions", you can include the code in this lpython notebook and also discuss the results in the writeup file.

Note: Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. In addition, Markdown cells can be edited by typically double-clicking the cell to enter edit mode.

Step 0: Load The Data

```
In [5]: # Load pickled data
import numpy as np
import csv
import time
import glob
import pickle
import matplotlib.pyplot as plt
import random
from numpy.random import rand
from sklearn.utils import shuffle
import tensorflow as tf
from tensorflow.contrib.layers import flatten

# TODO: Fill this in based on where you saved the training and testing data

training_file      = "./train.p"
validation_file    = "./valid.p"
testing_file       = "./test.p"

with open(training_file, mode='rb') as f:
    train = pickle.load(f)
with open(validation_file, mode='rb') as f:
    valid = pickle.load(f)
with open(testing_file, mode='rb') as f:
    test = pickle.load(f)

X_train, y_train = train['features'], train['labels']
X_valid, y_valid = valid['features'], valid['labels']
X_test, y_test = test['features'], test['labels']
```

Step 1: Dataset Summary & Exploration

The pickled data is a dictionary with 4 key/value pairs:

- 'features' is a 4D array containing raw pixel data of the traffic sign images, (num examples, width, height, channels).
- 'labels' is a 1D array containing the label/class id of the traffic sign. The file `signnames.csv` contains id -> name mappings for each id.
- 'sizes' is a list containing tuples, (width, height) representing the original width and height the image.
- 'coords' is a list containing tuples, (x1, y1, x2, y2) representing coordinates of a bounding box around the sign in the image. **THESE COORDINATES ASSUME THE ORIGINAL IMAGE. THE PICKLED DATA CONTAINS RESIZED VERSIONS (32 by 32) OF THESE IMAGES**

Complete the basic data summary below. Use python, numpy and/or pandas methods to calculate the data summary rather than hard coding the results. For example, the [pandas shape method](http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.shape.html) (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.shape.html>) might be useful for calculating some of the summary results.

Provide a Basic Summary of the Data Set Using Python, Numpy and/or

Pandas

```
In [6]: ### Replace each question mark with the appropriate value.  
### Use python, pandas or numpy methods rather than hard coding the results  
  
# TODO: Number of training examples  
n_train      = np.size(X_train, 0)  
  
# TODO: Number of validation examples  
n_validation = np.size(X_valid,0)  
  
# TODO: Number of testing examples.  
n_test       = np.size(X_test,0)  
  
# TODO: What's the shape of an traffic sign image?  
image_shape = X_train[0].shape  
  
# TODO: How many unique classes/labels there are in the dataset.  
n_classes = np.unique(y_train).size  
  
print("Number of training examples =", n_train)  
print("Number of testing examples =", n_test)  
print("Image data shape =", image_shape)  
print("Number of classes =", n_classes)  
  
Number of training examples = 34799  
Number of testing examples = 12630  
Image data shape = (32, 32, 3)  
Number of classes = 43
```

Include an exploratory visualization of the dataset

Visualize the German Traffic Signs Dataset using the pickled file(s). This is open ended, suggestions include: plotting traffic sign images, plotting the count of each sign, etc.

The [Matplotlib](http://matplotlib.org/) (<http://matplotlib.org/>) [examples](http://matplotlib.org/examples/index.html) (<http://matplotlib.org/examples/index.html>) and [gallery](http://matplotlib.org/gallery.html) (<http://matplotlib.org/gallery.html>) pages are a great resource for doing visualizations in Python.

NOTE: It's recommended you start with something simple first. If you wish to do more, come back to it after you've completed the rest of the sections. It can be interesting to look at the distribution of classes in the training, validation and test set. Is the distribution the same? Are there more examples of some classes than others?

```

In [7]: ### Data exploration visualization code goes here.
### Feel free to use as many code cells as needed.
# Visualizations will be shown in the notebook.
%matplotlib inline
index = random.randint(0, len(X_train))
image = X_train[index].squeeze()

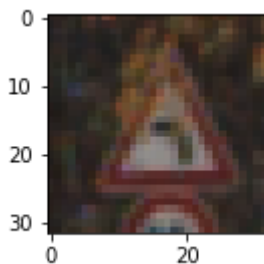
plt.figure(figsize = (2,2))
plt.imshow(image)
print(y_train[index])
#Reference
num_of_samples=[]
plt.figure(figsize=(12, 16.5))
for i in range(0, n_classes):
    plt.subplot(11, 4, i+1)
    x_selected = X_train[y_train == i]
    plt.imshow(x_selected[0, :, :, :]) #draw the first image of each class
    plt.title(i)
    plt.axis('off')
    num_of_samples.append(len(x_selected))
plt.show()

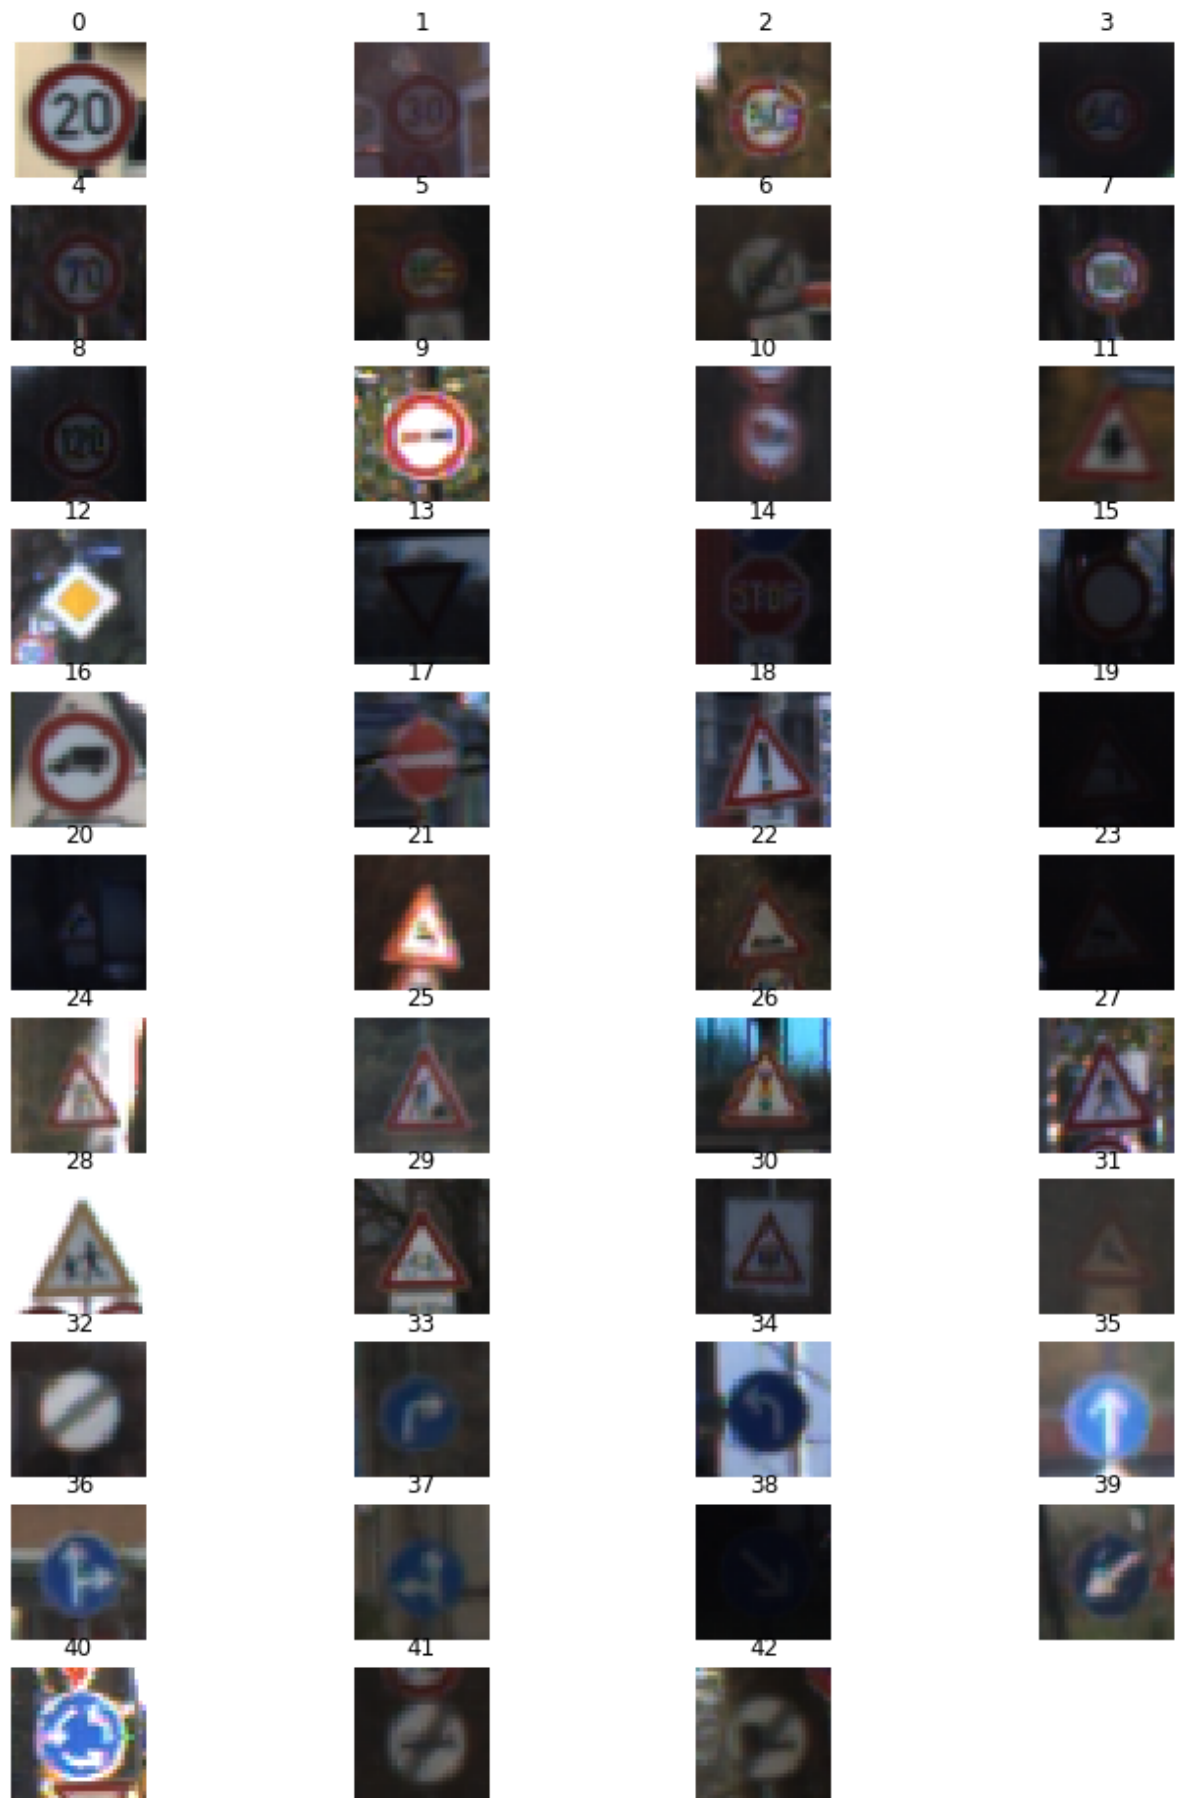
#Plot number of images per class
plt.figure(figsize=(12, 4))
plt.bar(range(0, n_classes), num_of_samples)
plt.title("Distribution of the train dataset")
plt.xlabel("Class number")
plt.ylabel("Number of images")
plt.show()

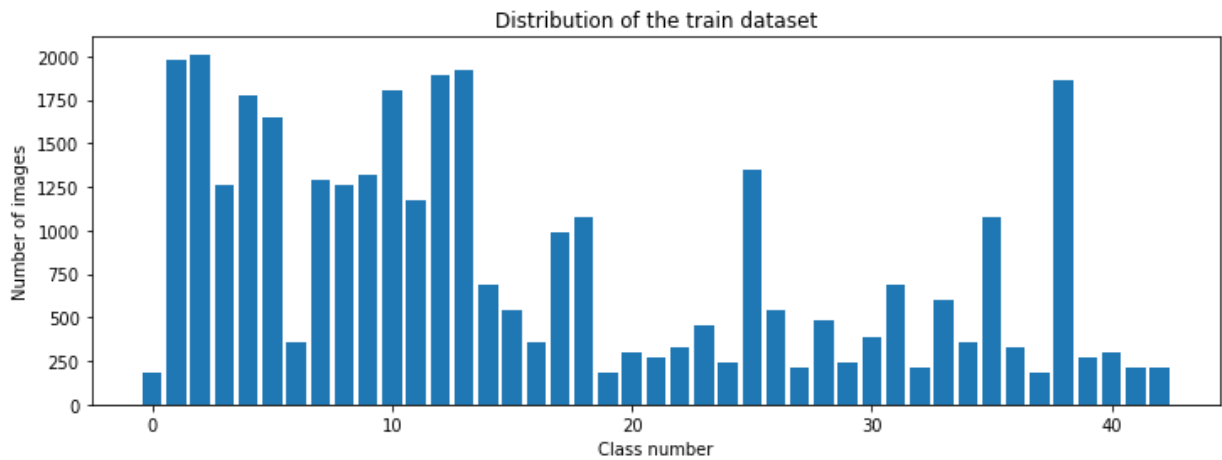
print("Min number of images per class =", min(num_of_samples))
print("Max number of images per class =", max(num_of_samples))

```

19







Min number of images per class = 180

Max number of images per class = 2010

Step 2: Design and Test a Model Architecture

Design and implement a deep learning model that learns to recognize traffic signs. Train and test your model on the [German Traffic Sign Dataset \(http://benchmark.ini.rub.de/?section=gtsrb&subsection=dataset\)](http://benchmark.ini.rub.de/?section=gtsrb&subsection=dataset).

The LeNet-5 implementation shown in the [classroom \(https://classroom.udacity.com/nanodegrees/nd013/parts/fbf77062-5703-404e-b60c-95b78b2f3f9e/modules/6df7ae49-c61c-4bb2-a23e-6527e69209ec/lessons/601ae704-1035-4287-8b11-e2c2716217ad/concepts/d4aca031-508f-4e0b-b493-e7b706120f81\)](https://classroom.udacity.com/nanodegrees/nd013/parts/fbf77062-5703-404e-b60c-95b78b2f3f9e/modules/6df7ae49-c61c-4bb2-a23e-6527e69209ec/lessons/601ae704-1035-4287-8b11-e2c2716217ad/concepts/d4aca031-508f-4e0b-b493-e7b706120f81) at the end of the CNN lesson is a solid starting point. You'll have to change the number of classes and possibly the preprocessing, but aside from that it's plug and play!

With the LeNet-5 solution from the lecture, you should expect a validation set accuracy of about 0.89. To meet specifications, the validation set accuracy will need to be at least 0.93. It is possible to get an even higher accuracy, but 0.93 is the minimum for a successful project submission.

There are various aspects to consider when thinking about this problem:

- Neural network architecture (is the network over or underfitting?)
- Play around preprocessing techniques (normalization, rgb to grayscale, etc)
- Number of examples per label (some have more than others).
- Generate fake data.

Here is an example of a [published baseline model on this problem \(http://yann.lecun.com/exdb/publis/pdf/sermanet-ijcnn-11.pdf\)](http://yann.lecun.com/exdb/publis/pdf/sermanet-ijcnn-11.pdf). It's not required to be familiar with the approach used in the paper but, it's good practice to try to read papers like these.

Pre-process the Data Set (normalization, grayscale, etc.)

Minimally, the image data should be normalized so that the data has mean zero and equal variance. For image data, $(\text{pixel} - 128) / 128$ is a quick way to approximately normalize the data and can

be used in this project.

Other pre-processing steps are optional. You can try different techniques to see if it improves performance.

Use the code cell (or multiple code cells, if necessary) to implement the first step of your project.

```
In [8]: ### Preprocess the data here. It is required to normalize the data. Other p
### converting to grayscale, etc.
### Feel free to use as many code cells as needed.
from sklearn.utils import shuffle

X_train, y_train = shuffle(X_train, y_train)

def rgb2gray(rgb):
    r, g, b = rgb[:, :, :, 0], rgb[:, :, :, 1], rgb[:, :, :, 2]
    gray = 0.2989 * r + 0.5870 * g + 0.1140 * b
    return gray

def normalize_grayscale(image_data):
    """
    Normalize the image data with Min-Max scaling to a range of [0.1, 0.9]
    :param image_data: The image data to be normalized
    :return: Normalized image data
    """
    a = 0.1
    b = 0.9
    grayscale_min = 0
    grayscale_max = 255
    return a + ( ( image_data - grayscale_min ) * ( b - a ) ) / ( grayscale_max -

X_normal = normalize_grayscale(rgb2gray(X_train))
X_normal = X_normal[..., np.newaxis]
print(X_train.shape)
print(X_normal.shape)
print()
X_valid_norm = normalize_grayscale(rgb2gray(X_valid))
X_valid_norm = X_valid_norm[..., np.newaxis]
print(X_valid.shape)
print(X_valid_norm.shape)
print()
X_test_norm = normalize_grayscale(rgb2gray(X_test))
X_test_norm = X_test_norm[..., np.newaxis]
print(X_test.shape)
print(X_test_norm.shape)

(34799, 32, 32, 3)
(34799, 32, 32, 1)

(4410, 32, 32, 3)
(4410, 32, 32, 1)

(12630, 32, 32, 3)
(12630, 32, 32, 1)
```

Model Architecture


```

In [12]: ### Define your architecture here.
### Feel free to use as many code cells as needed.
import tensorflow as tf
from tensorflow.contrib.layers import flatten
import matplotlib.pyplot as plt

EPOCHS      = 20
BATCH_SIZE  = 128
keep_prob   = tf.placeholder(tf.float32)

def LeNet(x):
    # Arguments used for tf.truncated_normal, randomly defines variables for
    mu = 0
    sigma = 0.1

    # SOLUTION: Layer 1: Convolutional. Input = 32x32x1. Output = 28x28x6.
    conv1_W = tf.Variable(tf.truncated_normal(shape=(5, 5, 1, 6), mean = mu,
    conv1_b = tf.Variable(tf.zeros(6))
    conv1    = tf.nn.conv2d(x, conv1_W, strides=[1, 1, 1, 1], padding='VALID')

    # SOLUTION: Activation.
    conv1 = tf.nn.relu(conv1)

    # SOLUTION: Pooling. Input = 28x28x6. Output = 14x14x6.
    conv1 = tf.nn.max_pool(conv1, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1],

    # SOLUTION: Layer 2: Convolutional. Output = 10x10x16.
    conv2_W = tf.Variable(tf.truncated_normal(shape=(5, 5, 6, 16), mean = mu,
    conv2_b = tf.Variable(tf.zeros(16))
    conv2    = tf.nn.conv2d(conv1, conv2_W, strides=[1, 1, 1, 1], padding='VA

    # SOLUTION: Activation.
    conv2 = tf.nn.relu(conv2)

    # SOLUTION: Pooling. Input = 10x10x16. Output = 5x5x16.
    conv2 = tf.nn.max_pool(conv2, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1],

    # SOLUTION: Flatten. Input = 5x5x16. Output = 400.
    fc0    = flatten(conv2)
    dpout_fc0 = tf.nn.dropout(fc0, keep_prob)

    # SOLUTION: Layer 3: Fully Connected. Input = 400. Output = 120.
    fc1_W = tf.Variable(tf.truncated_normal(shape=(400, 120), mean = mu, std
    fc1_b = tf.Variable(tf.zeros(120))

    fc1    = tf.matmul(dpout_fc0, fc1_W) + fc1_b

    # SOLUTION: Activation.
    fc1    = tf.nn.relu(fc1)

    # SOLUTION: Layer 4: Fully Connected. Input = 120. Output = 84.
    fc2_W = tf.Variable(tf.truncated_normal(shape=(120, 84), mean = mu, std
    fc2_b = tf.Variable(tf.zeros(84))
    #dpout_fc1 = tf.nn.dropout(fc1, keep_prob)

    fc2    = tf.matmul(fc1, fc2_W) + fc2_b

```

```

# SOLUTION: Activation.
fc2      = tf.nn.relu(fc2)

# SOLUTION: Layer 5: Fully Connected. Input = 84. Output = 43.
fc3_W    = tf.Variable(tf.truncated_normal(shape=(84, 43), mean = mu, stdc
fc3_b    = tf.Variable(tf.zeros(43))
logits   = tf.matmul(fc2, fc3_W) + fc3_b

return logits

#Features and Labels
#Train LeNet to classify data
#x is a placeholder for a batch of input images. y is a placeholder for a ba
x = tf.placeholder(tf.float32, (None, 32, 32, 1))
y = tf.placeholder(tf.int32, (None))
one_hot_y = tf.one_hot(y, 43)

```

Train, Validate and Test the Model

A validation set can be used to assess how well the model is performing. A low accuracy on the training and validation sets imply underfitting. A high accuracy on the training set but low accuracy on the validation set implies overfitting.

```

In [13]: ### Train your model here.
### Calculate and report the accuracy on the training and validation set.
### Once a final model architecture is selected,
### the accuracy on the test set should be calculated and reported as well.
### Feel free to use as many code cells as needed.
rate = 0.001
dropout = 0.5

logits = LeNet(x)
cross_entropy = tf.nn.softmax_cross_entropy_with_logits(labels=one_hot_y, logits=logits)
loss_operation = tf.reduce_mean(cross_entropy)
optimizer = tf.train.AdamOptimizer(learning_rate = rate)
training_operation = optimizer.minimize(loss_operation)

correct_prediction = tf.equal(tf.argmax(logits, 1), tf.argmax(one_hot_y, 1))
accuracy_operation = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
saver = tf.train.Saver()

def evaluate(X_data, y_data):
    num_examples = len(X_data)
    total_accuracy = 0
    sess = tf.get_default_session()
    for offset in range(0, num_examples, BATCH_SIZE):
        batch_x, batch_y = X_data[offset:offset+BATCH_SIZE], y_data[offset:offset+BATCH_SIZE]
        accuracy = sess.run(accuracy_operation, feed_dict={x: batch_x, y: batch_y})
        total_accuracy += (accuracy * len(batch_x))
    return total_accuracy / num_examples

#Running the training data through the training pipeline to train the model
with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())
    num_examples = len(X_normal)

    print("Training...")
    print()

    #Before each epoch, shuffle the training set.
    for i in range(EPOCHS):
        X_normal, y_train = shuffle(X_normal, y_train)
        for offset in range(0, num_examples, BATCH_SIZE):
            end = offset + BATCH_SIZE
            batch_x, batch_y = X_normal[offset:end], y_train[offset:end]
            sess.run(training_operation, feed_dict={x: batch_x, y: batch_y})

        #After each epoch, measure the loss and accuracy of the validation set.
        validation_accuracy = evaluate(X_valid_norm, y_valid)
        print("EPOCH {} ...".format(i+1))
        print("Validation Accuracy = {:.3f}".format(validation_accuracy))
        print()

    #Save the model after training
    saver.save(sess, './lenet')
    print("Model saved")

```

Training...

EPOCH 1 ...

Validation Accuracy = 0.497

```
EPOCH 2 ...  
Validation Accuracy = 0.780  
  
EPOCH 3 ...  
Validation Accuracy = 0.844  
  
EPOCH 4 ...  
Validation Accuracy = 0.880  
  
EPOCH 5 ...  
Validation Accuracy = 0.893  
  
EPOCH 6 ...  
Validation Accuracy = 0.903  
  
EPOCH 7 ...  
Validation Accuracy = 0.900  
  
EPOCH 8 ...  
Validation Accuracy = 0.919  
  
EPOCH 9 ...  
Validation Accuracy = 0.918  
  
EPOCH 10 ...  
Validation Accuracy = 0.927  
  
EPOCH 11 ...  
Validation Accuracy = 0.929  
  
EPOCH 12 ...  
Validation Accuracy = 0.925  
  
EPOCH 13 ...  
Validation Accuracy = 0.927  
  
EPOCH 14 ...  
Validation Accuracy = 0.944  
  
EPOCH 15 ...  
Validation Accuracy = 0.926  
  
EPOCH 16 ...  
Validation Accuracy = 0.928  
  
EPOCH 17 ...  
Validation Accuracy = 0.938  
  
EPOCH 18 ...  
Validation Accuracy = 0.932  
  
EPOCH 19 ...  
Validation Accuracy = 0.935  
  
EPOCH 20 ...  
Validation Accuracy = 0.942
```

Model saved

Step 3: Test a Model on New Images

To give yourself more insight into how your model is working, download at least five pictures of German traffic signs from the web and use your model to predict the traffic sign type.

You may find `signnames.csv` useful as it contains mappings from the class id (integer) to the actual sign name.

Load and Output the Images

```

In [14]: ### Load the images and plot them here.
### Feel free to use as many code cells as needed.
#reading in an image
#Import traffic signs class names
import glob
import csv
from PIL import Image

def rgb2gray(rgb):
    r, g, b = rgb[:, :, :, 0], rgb[:, :, :, 1], rgb[:, :, :, 2]
    gray = 0.2989 * r + 0.5870 * g + 0.1140 * b
    return gray

signs_img_class = []
with open('signnames.csv', 'rt') as csvfile:
    reader = csv.DictReader(csvfile, delimiter=',')
    for row in reader:
        signs_img_class.append((row['SignName']))

#Import test images

filelist = glob.glob('new_images/*.jpg')
test_img = np.array([np.array(Image.open(fname)) for fname in filelist])
test_img_gray = rgb2gray(test_img)
test_img_gray_mod = test_img_gray[..., np.newaxis]
#Visualize new raw images
plt.figure(figsize=(6, 6))
for i in range(8):
    plt.subplot(2, 4, i+1)
    plt.imshow(test_img[i])
    plt.title(i)
    plt.axis('on')
    num_of_samples.append(len(x_selected))
plt.show()
print(test_img_gray_mod.shape)

```



(8, 32, 32, 1)

Predict the Sign Type for Each Image

```
In [15]: ### Run the predictions here and use the model to output the prediction for
### Make sure to pre-process the images with the same pre-processing pipeline
### Feel free to use as many code cells as needed.
def test_lenet(X_data, sess):
    pred_sign = sess.run(tf.argmax(logits, 1), feed_dict={x: X_data, keep_prob: 1.0})
    return pred_sign

#Run Testing
with tf.Session() as sess:
    saver.restore(sess, './lenet')
    signs_classes = test_lenet(test_img_gray_mod, sess)

plt.figure(figsize=(12, 8))
for i in range(8):
    plt.subplot(2, 4, i+1)
    plt.imshow(test_img[i])
    plt.title(signs_img_class[signs_classes[i]])
    plt.axis('off')
plt.show()
```

Priority road



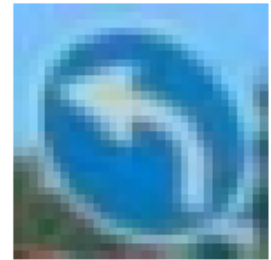
Right-of-way at the next intersection



General caution



Turn left ahead



Speed limit (60km/h)



Road work



Speed limit (30km/h)



Keep right



Analyze Performance

```

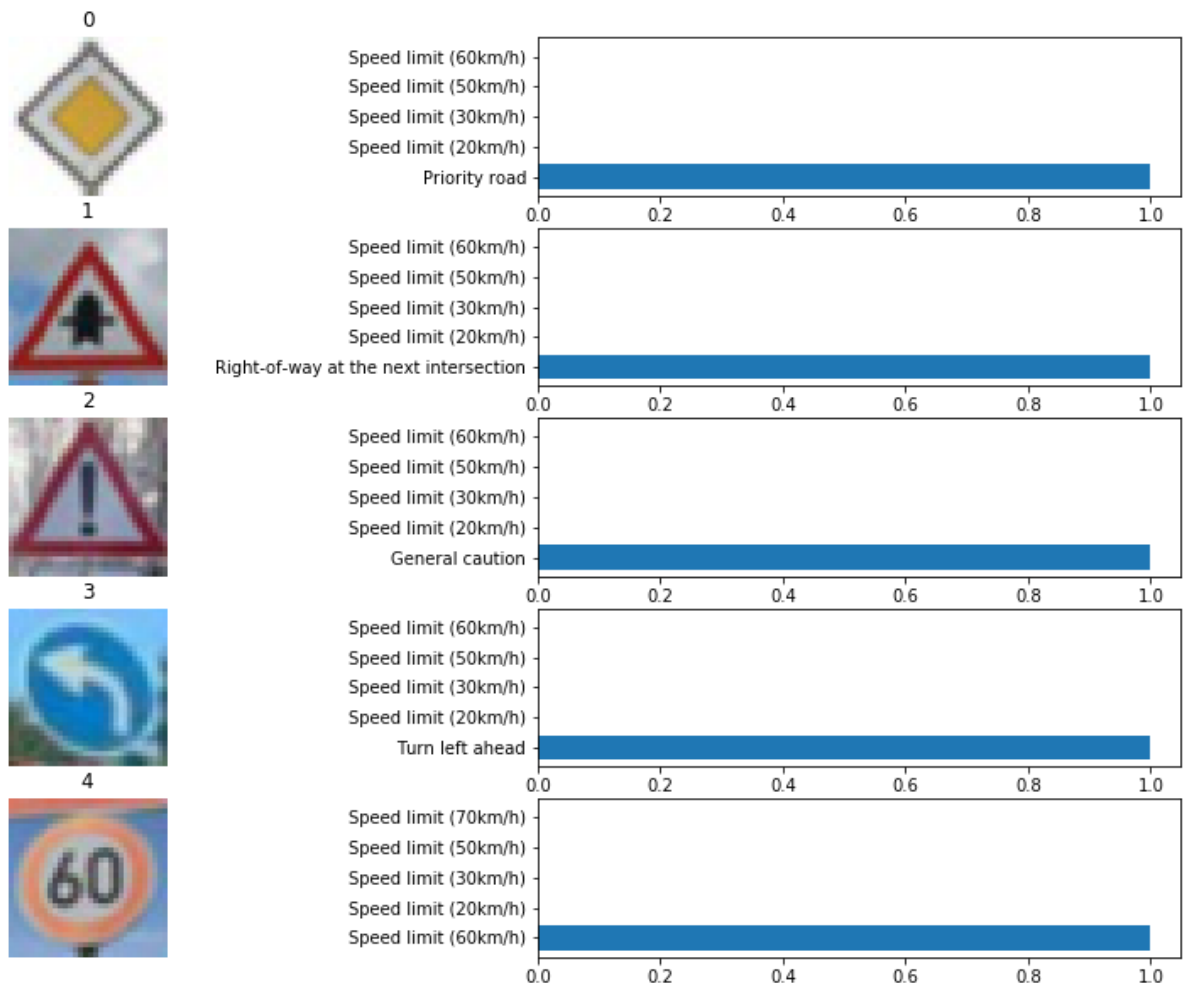
In [16]: ### Calculate the accuracy for these 5 new images.
### For example, if the model predicted 1 out of 5 signs correctly, it's 20%

def testsys(X_data, sess):
    prob = sess.run(tf.nn.softmax(logits), feed_dict={x: X_data, keep_prob:
    top_5 = tf.nn.top_k(prob, k = 5)
    return sess.run(top_5)

with tf.Session() as sess:
    saver.restore(sess, './lenet')
    signs_top_5 = testsys(test_img_gray_mod, sess)

plt.figure(figsize = (15, 10))
for i in range (5):
    plt.subplot(5, 2, 2*i+1)
    plt.imshow(test_img[i])
    plt.title(i)
    plt.axis('off')
    plt.subplot(5, 2, 2*i+2)
    plt.barh(np.arange(1, 6, 1), signs_top_5.values[i, :])
    labels = [signs_img_class[j] for j in signs_top_5.indices[i]]
    plt.yticks(np.arange(1, 6, 1), labels)
plt.show()

```



Output Top 5 Softmax Probabilities For Each Image Found on the Web

For each of the new images, print out the model's softmax probabilities to show the **certainty** of the model's predictions (limit the output to the top 5 probabilities for each image). `tf.nn.top_k` (https://www.tensorflow.org/versions/r0.12/api_docs/python/nntop_k) could prove helpful here.

The example below demonstrates how `tf.nn.top_k` can be used to find the top k predictions for each image.

`tf.nn.top_k` will return the values and indices (class ids) of the top k predictions. So if `k=3`, for each sign, it'll return the 3 largest probabilities (out of a possible 43) and the corresponding class ids.

Take this numpy array as an example. The values in the array represent predictions. The array contains softmax probabilities for five candidate images with six possible classes. `tf.nn.top_k` is used to choose the three classes with the highest probability:

```
# (5, 6) array
a = np.array([[ 0.24879643,  0.07032244,  0.12641572,  0.34763842,
                0.07893497,
                0.12789202],
              [ 0.28086119,  0.27569815,  0.08594638,  0.0178669 ,  0.18063
401,
                0.15899337],
              [ 0.26076848,  0.23664738,  0.08020603,  0.07001922,  0.11343
71 ,
                0.23892179],
              [ 0.11943333,  0.29198961,  0.02605103,  0.26234032,  0.13513
48 ,
                0.16505091],
              [ 0.09561176,  0.34396535,  0.0643941 ,  0.16240774,  0.24206
137,
                0.09155967]])
```

Running it through `sess.run(tf.nn.top_k(tf.constant(a), k=3))` produces:

```
TopKV2(values=array([[ 0.34763842,  0.24879643,  0.12789202],
                    [ 0.28086119,  0.27569815,  0.18063401],
                    [ 0.26076848,  0.23892179,  0.23664738],
                    [ 0.29198961,  0.26234032,  0.16505091],
                    [ 0.34396535,  0.24206137,  0.16240774]]), indices=array([[3,
0, 5],
                    [0, 1, 4],
                    [0, 5, 1],
                    [1, 3, 5],
                    [1, 4, 3]], dtype=int32))
```

Looking just at the first row we get `[0.34763842, 0.24879643, 0.12789202]`, you can confirm these are the 3 largest probabilities in `a`. You'll also notice `[3, 0, 5]` are the corresponding indices.

```
In [ ]: ### Print out the top five softmax probabilities for the predictions on the  
### Feel free to use as many code cells as needed.
```

Project Writeup

Once you have completed the code implementation, document your results in a project writeup using this [template \(https://github.com/udacity/CarND-Traffic-Sign-Classifer-Project/blob/master/writeup_template.md\)](https://github.com/udacity/CarND-Traffic-Sign-Classifer-Project/blob/master/writeup_template.md) as a guide. The writeup can be in a markdown or pdf file.

Note: Once you have completed all of the code implementations and successfully answered each question above, you may finalize your work by exporting the iPython Notebook as an HTML document. You can do this by using the menu above and navigating to "\n", "**File -> Download as -> HTML (.html)**". Include the finished document along with this notebook as your submission.

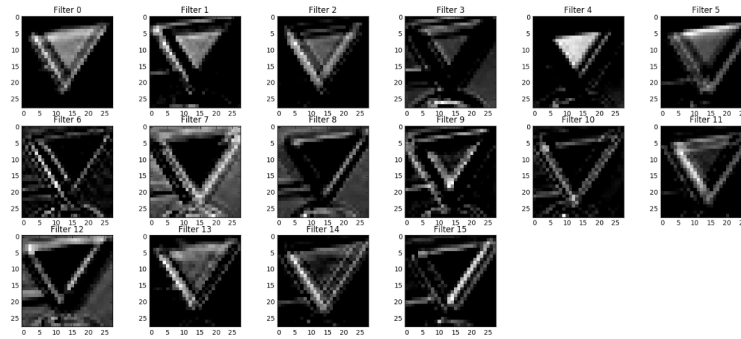
Step 4 (Optional): Visualize the Neural Network's State with Test Images

This Section is not required to complete but acts as an additional exercise for understanding the output of a neural network's weights. While neural networks can be a great learning device they are often referred to as a black box. We can understand what the weights of a neural network look like better by plotting their feature maps. After successfully training your neural network you can see what its feature maps look like by plotting the output of the network's weight layers in response to a test stimuli image. From these plotted feature maps, it's possible to see what characteristics of an image the network finds interesting. For a sign, maybe the inner network feature maps react with high activation to the sign's boundary outline or to the contrast in the sign's painted symbol.

Provided for you below is the function code that allows you to get the visualization output of any tensorflow weight layer you want. The inputs to the function should be a stimuli image, one used during training or a new one you provided, and then the tensorflow variable name that represents the layer's state during the training process, for instance if you wanted to see what the [LeNet lab's \(https://classroom.udacity.com/nanodegrees/nd013/parts/fbf77062-5703-404e-b60c-95b78b2f3f9e/modules/6df7ae49-c61c-4bb2-a23e-6527e69209ec/lessons/601ae704-1035-4287-8b11-e2c2716217ad/concepts/d4aca031-508f-4e0b-b493-e7b706120f81\)](https://classroom.udacity.com/nanodegrees/nd013/parts/fbf77062-5703-404e-b60c-95b78b2f3f9e/modules/6df7ae49-c61c-4bb2-a23e-6527e69209ec/lessons/601ae704-1035-4287-8b11-e2c2716217ad/concepts/d4aca031-508f-4e0b-b493-e7b706120f81) feature maps looked like for its second convolutional layer you could enter conv2 as the tf_activation variable.

For an example of what feature map outputs look like, check out NVIDIA's results in their paper [End-to-End Deep Learning for Self-Driving Cars \(https://devblogs.nvidia.com/parallelforall/deep-learning-self-driving-cars/\)](https://devblogs.nvidia.com/parallelforall/deep-learning-self-driving-cars/) in the section Visualization of internal CNN State. NVIDIA was able to show that their network's inner weights had high activations to road boundary lines by comparing feature maps from an image with a clear path to one without. Try experimenting with a similar test to show

that your trained network's weights are looking for interesting features, whether it's looking at differences in feature maps from images with or without a sign, or even what feature maps look like in a trained network vs a completely untrained one on the same sign image.



Your output should look something like this (above)

```
In [ ]: # Visualize your network's feature maps here.
# Feel free to use as many code cells as needed.

image_input: the test image being fed into the network to produce the feature
tf_activation: should be a tf variable name used during your training procedu
activation_min/max: can be used to view the activation contrast in more detai
plt_num: used to plot out multiple different weight feature map sets on the s

outputFeatureMap(image_input, tf_activation, activation_min=-1, activation_
# Here make sure to preprocess your image_input in a way your network expec
# with size, normalization, ect if needed
image_input = X_normal[index].squeeze()
# Note: x should be the same name as your network's tensorflow data placeho
# If you get an error tf_activation is not defined it may be having trouble
activation = tf_activation.eval(session=sess, feed_dict={x : image_input})
featuremaps = activation.shape[3]
plt.figure(plt_num, figsize=(15,15))
for featuremap in range(featuremaps):
    plt.subplot(6,8, featuremap+1) # sets the number of feature maps to sho
    plt.title('FeatureMap ' + str(featuremap)) # displays the feature map n
    if activation_min != -1 & activation_max != -1:
        plt.imshow(activation[0,::, featuremap], interpolation="nearest",
    elif activation_max != -1:
        plt.imshow(activation[0,::, featuremap], interpolation="nearest",
    elif activation_min != -1:
        plt.imshow(activation[0,::, featuremap], interpolation="nearest",
    else:
        plt.imshow(activation[0,::, featuremap], interpolation="nearest",
```

```
In [ ]: with tf.Session() as sess:
    saver.restore(sess, './lenet')
    print("conv1 : First layer")
    outputFeatureMap(test_img_gray_mod, conv1)
```

In []: